

# Fast Convex Optimization for Two-Layer ReLU Networks:

Equivalent Model Classes and Cone Decompositions

Aaron Mishkin

`amishkin@cs.stanford.edu`



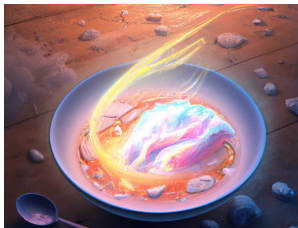
# Motivation: Neural Networks

---

Neural networks are building blocks for modern learning systems.

- **Vision:** object recognition, localization, and bounding.
- **Language:** audio transcription and machine translation.
- **Control:** robotics, autonomous cars, game-playing, ...

*A bowl of soup that is a portal to another dimension as digital art.*



Generated by DALL·E 2

But optimizing neural networks is **hard**!


## Motivation: Non-Convex Optimization

---

DALL·E 2 has 5.5 billion parameters and took billions of iterations to fit [2].

Neural network optimization is **non-convex**!


- **NP-Hard**: many sub-optimal local minima, saddles.
- **Speed/Stability**: tuning is critical for performance.
- **Model Churn**: hyper-parameters affect the final model [1].



# Convex Reformulations

## Non-Convex Problem


$$\min_{w, \alpha} \left\| \sum_{j=1}^m (Xw_j)_+ \alpha_j - y \right\|_2^2 + \lambda \sum_{j=1}^m \|w_j\|_2^2 + |\alpha_j|^2$$



# Convex Reformulations

## Non-Convex Problem


$$\min_{w, \alpha} \left\| \sum_{j=1}^m (Xw_j)_+ \alpha_j - y \right\|_2^2 + \lambda \sum_{j=1}^m \|w_j\|_2^2 + |\alpha_j|^2$$



## Convex Reformulation

$$\min_u \left\| \sum_{j=1}^p D_j X u_j - y \right\|_2^2 + \lambda \sum_{j=1}^p \|u_j\|_2^2,$$

where  $D_j = \text{diag}[\mathbb{1}(Xg_i \geq 0)]$



# Convex Reformulations: A Huge-Scale Linear Model

**Convex Form :** 
$$\min_u \left\| \sum_{j=1}^p D_j X u_j - y \right\|_2^2 + \lambda \sum_{j=1}^p \|u_j\|_2,$$

where  $D_j = \text{diag}[\mathbb{1}(X g_i \geq 0)]$

- **Exponential in general:**  $p \in O(r \cdot (\frac{n}{r})^r)$ , where  $r = \text{rank}(X)$ .
- But,  $D_j X$  is **row-sparse**,  $u$  is **sparse** when  $\lambda \gg 0$ , and the objective is **quadratic** + simple non-smooth.

# Convex Reformulations: A Huge-Scale Linear Model

**Convex Form :** 
$$\min_u \left\| \sum_{j=1}^p D_j X u_j - y \right\|_2^2 + \lambda \sum_{j=1}^p \|u_j\|_2,$$

where  $D_j = \text{diag}[\mathbb{1}(X g_i \geq 0)]$

- **Exponential in general:**  $p \in O(r \cdot (\frac{n}{r})^r)$ , where  $r = \text{rank}(X)$ .
- But,  $D_j X$  is **row-sparse**,  $u$  is **sparse** when  $\lambda \gg 0$ , and the objective is **quadratic** + simple non-smooth.

---

Solve with (accelerated) **proximal-gradient** methods:

$$u_i^+ = u_i^k - \eta_k X^\top D_i^\top \left( \sum_{j=1}^p D_j X u_j - y \right)$$
$$u_i^{k+1} = u_i^+ \left( 1 - \frac{\eta_k \cdot \lambda}{\|u_i^+\|_2} \right)_+$$

## Convex Reformulations: Performance

---

Fast solvers use **numerical tricks** and **hardware acceleration**:

- **Classic Tricks**: faster convergence via line-search, restarts, ...
- **CUDA GPUs**: 70× faster Mat-Vec operations using float32.
- **Code Optimization**: tensor operations remove intermediate computations, data normalization improves conditioning, ...


**Scaling** is still a problem!

- Dense operations on GPUs are faster than sparse computations with OpenMP — does this reverse at scale?
- GPU memory is typically 32GB, but ImageNet is 150GB — can we use multi-GPU programming models?
- How can we leverage the (dynamic) sparsity pattern of  $u^k$ ?




Thanks for Listening!

## Bonus: Numerical Results (1)



## Bonus: Numerical Results (2)



# References I

---



Peter Henderson et al. “Deep Reinforcement Learning That Matters”. In: *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence, (AAAI-18), New Orleans, Louisiana, USA, February 2-7, 2018*. 2018, pp. 3207–3214.



Aditya Ramesh et al. “Hierarchical Text-Conditional Image Generation with CLIP Latents”. In: *CoRR* abs/2204.06125 (2022). arXiv: 2204.06125.